A decorative graphic on the left side of the slide consisting of two overlapping parallelograms. The front one is blue and the back one is a light greenish-blue. They are positioned diagonally, with the blue one partially covering the green one.

CPU-Based Shortest Path Algorithms

Brendan Ngo, Aravind Srinivasan




Definition: Shortest Path

Given a digraph G with non-negative weights on its edges, the shortest path from a source node S , and a destination node T is a directed path from S to T with the minimum total weight.

- Single-Source Shortest Path: S to V_i , $\forall V_i \in V$
- All-Pairs Shortest Path: V_i to V_j , $\forall V_i, V_j \in V$

Example



	A	B	C	D
A	0	1	2	4
B	∞	0	1	3
C	∞	∞	0	∞
D	∞	∞	∞	0



SSSP: Single-Source Shortest Path

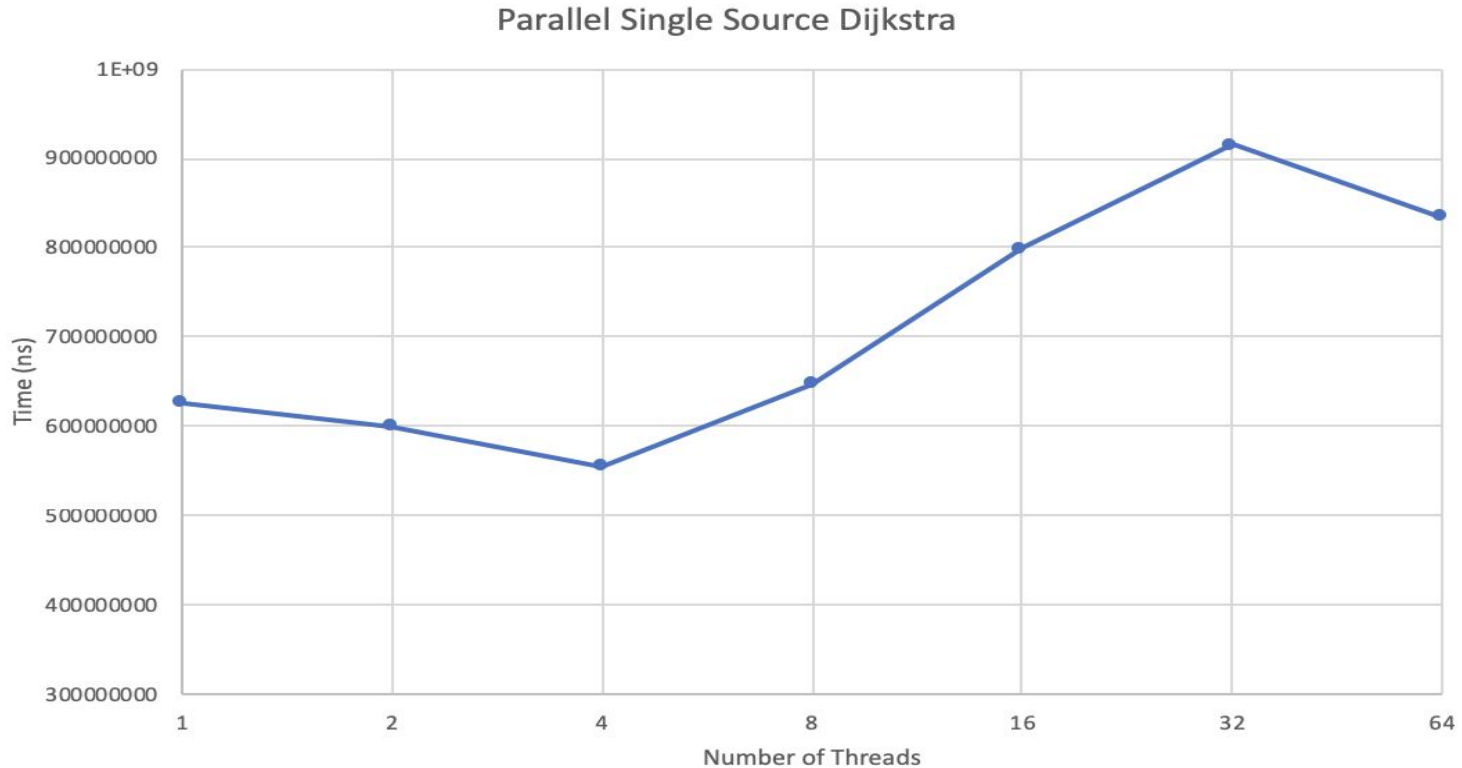
1. Mark all nodes as unvisited and set their weights to ∞
2. Set current node weight to 0.
3. Remove node with smallest weight from unvisited set
4. Relax its outgoing edges - Given edge (u, v) , $\text{dist}(v) = \min(\text{dist}(u) + \text{weight}(e), \text{dist}(v))$
5. Repeat steps 3 and 4 until the unvisited set is empty.

Dijkstra's Runtime: $O(E + V \lg V)$

Parallel Dijkstra's finds the minimum in parallel with P threads.

Runtime: $O(E + V^2/P) \rightarrow O(E + V)$

SSSP: Single-Source Shortest Path





APSP: All-Pairs Shortest Path

Sequential Floyd-Warshall Algorithm $O(V^3)$:


```
1 let dist be a  $|V| \times |V|$  array of minimum distances initialized to  $\infty$  (infinity)
2 for each edge  $(u,v)$ 
3      $\text{dist}[u][v] \leftarrow w(u,v)$  // the weight of the edge  $(u,v)$ 
4 for each vertex  $v$ 
5      $\text{dist}[v][v] \leftarrow 0$ 
6 for  $k$  from 1 to  $|V|$ 
7     for  $i$  from 1 to  $|V|$ 
8         for  $j$  from 1 to  $|V|$ 
9             if  $\text{dist}[i][j] > \text{dist}[i][k] + \text{dist}[k][j]$ 
10                  $\text{dist}[i][j] \leftarrow \text{dist}[i][k] + \text{dist}[k][j]$ 
11             end if
```

Lines 7-11 can be parallelized with P threads for runtime $O(V^3/P) \rightarrow O(V)$ with V^2 threads

APSP: All-Pairs Shortest Path

Alternative: Run Dijkstra's Algorithm in parallel $\forall V_i \in V$

Runtime: $O(E + V \lg V)$



As expected, Floyd-Warshall's is faster than running Parallel Dijkstra's


Trade off: Floyd-Warshall needs V^2 threads vs Dijkstra's needs V threads.

When we have large, sparse graphs with limited cores,

$$O(V/P * (E + V \lg V)) \approx O(V^2 \lg V / P) \leq O(V^3 / P)$$

Matrix Multiplication Based Implementations: Distance Products

Given W , a $n \times n$ matrix containing the edge weights of the graph, W^k gives the distances realized by paths that use at most k edges.



1.

$$(\mathbf{A} \star \mathbf{B})_{ij} = \min_{k=1}^n \{a_{ik} + b_{kj}\}, 1 \leq i, j \leq n.$$

2.

$$\mathbf{A} = \begin{pmatrix} 0 & 1 & 0 & 0 & 0 & 1 \\ 1 & 0 & 1 & 1 & 1 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 & 0 & 1 \\ 1 & 0 & 0 & 0 & 1 & 0 \end{pmatrix}$$

$$\mathbf{A}^2 = \begin{pmatrix} 2 & 0 & 1 & 1 & 2 & 0 \\ 0 & 4 & 1 & 2 & 1 & 2 \\ 1 & 1 & 2 & 1 & 2 & 0 \\ 1 & 2 & 1 & 3 & 1 & 1 \\ 2 & 1 & 2 & 1 & 3 & 0 \\ 0 & 2 & 0 & 1 & 0 & 2 \end{pmatrix}$$

$$\mathbf{B} = \begin{pmatrix} 0 & 1 & 1 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 & 1 & 1 \\ 1 & 1 & 0 & 1 & 1 & 0 \\ 1 & 1 & 1 & 0 & 1 & 1 \\ 1 & 1 & 1 & 1 & 0 & 1 \\ 1 & 1 & 0 & 1 & 1 & 0 \end{pmatrix}$$

Seidel's Algorithm

- APSP Algorithm for unweighted, undirected graphs
- Computes logarithmic number of distance matrices
- $O(M(n) \cdot \log(n))$, where $M(n)$ is the cost of multiplying 2 matrices
 - Sequential: $O(n^{2.375477} \log(n))$
 - Parallel: $O(\log^2(n))$

Seidel(A)

Let $Z = A \cdot A$

Let B be an $n \times n$ 0-1 matrix,

where $b_{ij} = \begin{cases} 1 & \text{if } i \neq j \text{ and } (a_{ij} = 1 \text{ or } z_{ij} > 0) \\ 0 & \text{otherwise} \end{cases}$

IF $\forall i, j, i \neq j, b_{ij} = 1$ then

Return $D = 2B - A$

Let $T = \text{Seidel}(B)$

Let $X = T \cdot A$

Return $n \times n$ matrix D ,

where $d_{ij} = \begin{cases} 2t_{ij} & \text{if } x_{ij} \geq t_{ij} \cdot \text{degree}(j) \\ 2t_{ij} - 1 & \text{if } x_{ij} < t_{ij} \cdot \text{degree}(j) \end{cases}$



Shoshan-Zwick Algorithm

- APSP algorithm that works on weighted, undirected graphs where existing edge weights are integers in the range of $\{1, 2, \dots, N\}$
- Computes a logarithmic number of distance products in order to determine shortest paths
- Algorithm is based on not allowing the range of elements in the matrix it is uses to increase
- Stores distance matrices as it iterates in order to recover the shortest path
- $O(\log(N*n)*N*M(n))$
 - Sequential: $O(N*n^{2.375477}*\log(N*n))$
 - Parallel: $O(N*\log(n)\log^2N*n)$

Function SHOSHAN-ZWICK-APSP(D)

```

1:  $l = \lceil \log_2 n \rceil$ .
2:  $m = \log_2 M$ .
3: for ( $k = 1$  to  $m + 1$ ) do
4:    $\mathbf{D} = \text{clip}(\mathbf{D} \star \mathbf{D}, 0, 2 \cdot M)$ .
5: end for
6:  $\mathbf{A}_0 = \mathbf{D} - M$ .
7: for ( $k = 1$  to  $l$ ) do
8:    $\mathbf{A}_k = \text{clip}(\mathbf{A}_{k-1} \star \mathbf{A}_{k-1}, -M, M)$ .
9: end for
10:  $\mathbf{C}_l = -M$ .
11:  $\mathbf{P}_l = \text{clip}(\mathbf{D}, 0, M)$ .
12:  $\mathbf{Q}_l = +\infty$ .
13: for ( $k = l - 1$  down to 0) do
14:    $\mathbf{C}_k = [\text{clip}(\mathbf{P}_{k+1} \star \mathbf{A}_k, -M, M) \wedge \mathbf{C}_{k+1}] \vee [\text{clip}(\mathbf{Q}_{k+1} \star \mathbf{A}_k, -M, M) \bar{\wedge} \mathbf{C}_{k+1}]$ .
15:    $\mathbf{P}_k = \mathbf{P}_{k+1} \vee \mathbf{Q}_{k+1}$ .
16:    $\mathbf{Q}_k = \text{chop}(\mathbf{C}_k, 1 - M, M)$ .
17: end for
18: for ( $k = 1$  to  $l$ ) do
19:    $\mathbf{B}_k = (\mathbf{C}_k \geq 0)$ .
20: end for
21:  $\hat{\mathbf{B}}_0 = (-M < \mathbf{P}_0 < 0)$ .
22:  $\hat{\mathbf{R}} = \mathbf{P}_0$ .
23:  $\Delta = M \cdot \sum_{k=1}^l 2^k \cdot \mathbf{B}_k + 2 \cdot M \cdot \hat{\mathbf{B}}_0 + \hat{\mathbf{R}}$ .
24: return  $\Delta$ .

```

$$(\text{clip}(\mathbf{A}, a, b))_{ij} = \begin{cases} a & \text{if } a_{ij} < a \\ a_{ij} & \text{if } a \leq a_{ij} \leq b \\ +\infty & \text{if } a_{ij} > b \end{cases}$$

$$(\text{chop}(\mathbf{A}, a, b))_{ij} = \begin{cases} a_{ij} & \text{if } a \leq a_{ij} \leq b \\ +\infty & \text{otherwise} \end{cases}$$


$$(\mathbf{A} \wedge \mathbf{B})_{ij} = \begin{cases} a_{ij} & \text{if } b_{ij} < 0 \\ +\infty & \text{otherwise} \end{cases}$$

$$(\mathbf{A} \bar{\wedge} \mathbf{B})_{ij} = \begin{cases} a_{ij} & \text{if } b_{ij} \geq 0 \\ +\infty & \text{otherwise} \end{cases}$$

$$(\mathbf{A} \vee \mathbf{B})_{ij} = \begin{cases} a_{ij} & \text{if } a_{ij} \neq +\infty \\ b_{ij} & \text{if } a_{ij} = +\infty, b_{ij} \neq +\infty \\ +\infty & \text{if } a_{ij} = b_{ij} = +\infty \end{cases}$$

$$(\mathbf{C} \geq 0)_{ij} = \begin{cases} 1 & \text{if } c_{ij} \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

$$(0 \leq \mathbf{P} \leq M)_{ij} = \begin{cases} 1 & \text{if } 0 \leq p_{ij} \leq M \\ 0 & \text{otherwise} \end{cases}$$



Analysis of Matrix Based Implementations




Conclusion

- Shortest Path Algorithms can be improved using parallelization on CPU
- CPU-based implementations efficiency capped by number of available processors
 - Maximum speedup when num of threads = number of available processors (4)
- Speedups are only realized for larger number of nodes
 - For smaller inputs, benefits of parallelization are negligible and cost of setting up parallelization too high
- Matrix-based algorithms do not perform well in practice on CPU



References

- [1] Dijkstra's Algorithm, Dijkstra, E. W. (1959). "A note on two problems in connexion with graphs"
- [2] Floyd-Warshall Algorithm, Floyd, Robert W. (June 1962). "Algorithm 97: Shortest Path". *Communications of the ACM*. 5 (6): 345.
- [3] Seidel's Algorithm, Seidel, R. (1995). "On the All-Pairs-Shortest-Path Problem in Unweighted Undirected Graphs". *Journal of Computer and System Sciences*. **51** (3): 400–403.
- [4] Strassen Matrix Multiplication Algorithm, Strassen, Volker, *Gaussian Elimination is not Optimal*, Numer. Math. 13, p. 354-356, 1969
- [5] Coppersmith-Winograd Matrix Multiplication Algorithm, Coppersmith, Don; Winograd, Shmuel (1990), "Matrix multiplication via arithmetic progressions" (PDF), *Journal of Symbolic Computation*, **9** (3): 251
- [7] Shoshan-Zwick Algorithm, A. Shoshan and U. Zwick. All pairs shortest paths in undirected graphs with integer weights. In FOCS, pages 605–614, 1999.
- [8] Correction of the Shoshan-Zwick Algorithm, Pavlos Eirinakis, Matthew Williamson, and K. Subramani. On the Shoshan-Zwick algorithm for the all-pairs shortest path problem. *J. Graph Algorithms Appl.*, 21(2):177–181, 2017.



Parallelization is not helpful for smaller graphs